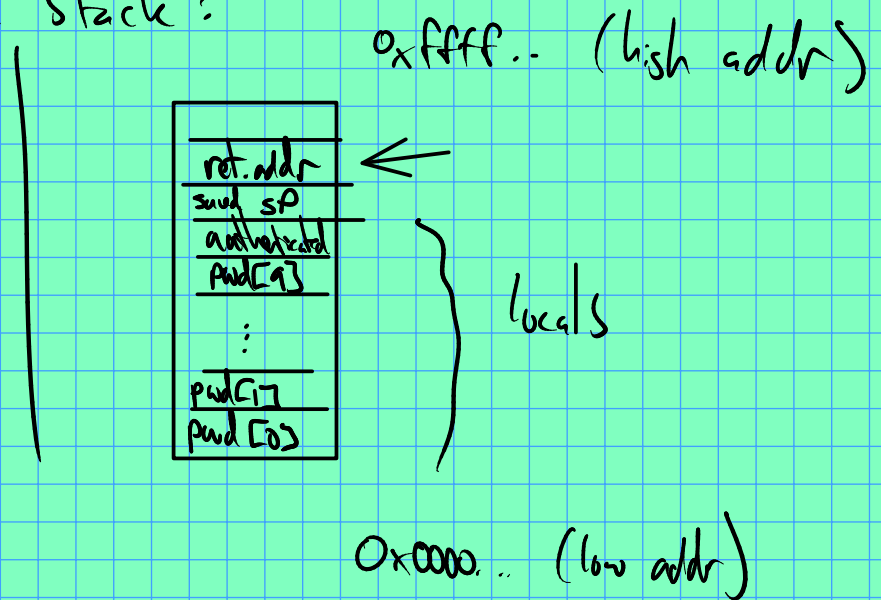


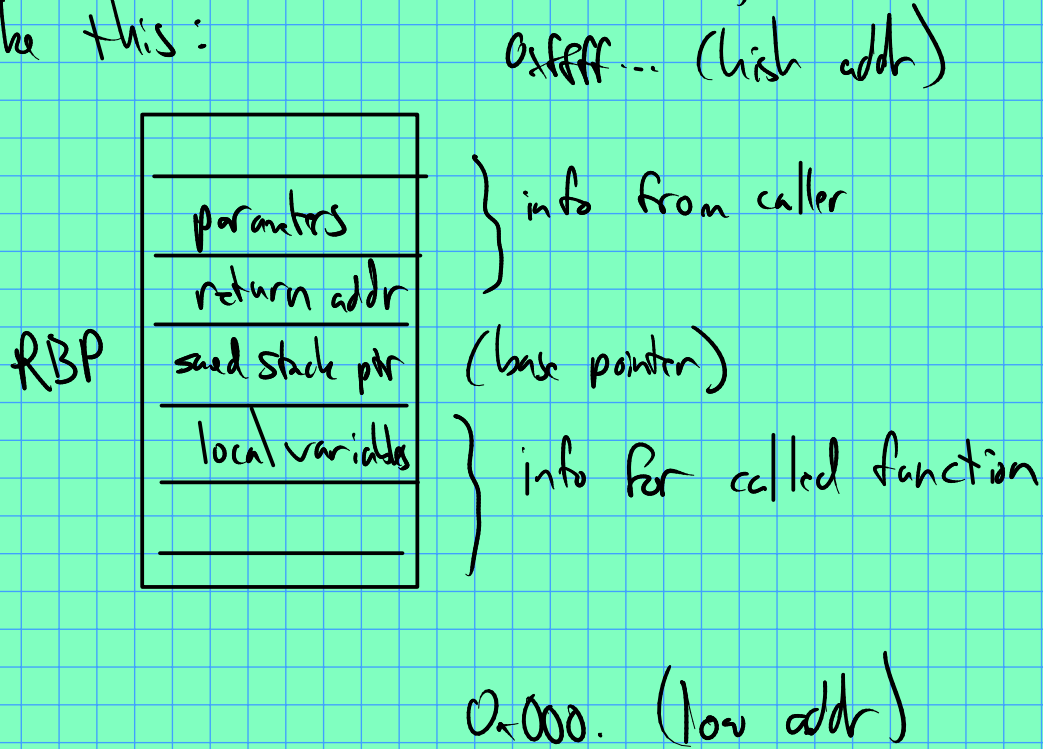
Software security: control-flow integrity.

local variables on call stack:

```
int authenticated = 0;  
char pwd[10];
```



Quick review on call stack. For x86_64, looks something like this:




Observations: call stack is fragile... by manipulating return address, we can change what instructions a program performs!

old school "stack smashing"

- Setup:
- ① victim program uses a fixed length array on its stack to hold user input.
 - ② user input is read into array with no bounds checks (e.g. via `scanf("%s")`)

prepare attack string like this:



Hurdles for attacker:

proper return address seems hard to guess...
however, one can improve the odds of success using
a "NOP-slide"

NOP NOP NOP NOP <Nasty code....>

if we ~~guess~~ anything in this range, we win!

(Note: NOP can be useful, e.g. for aligning code
on "nice" (16N) byte boundaries...)

Other issues: might have to be creative to avoid \x00
(null characters) in your attack code!

What can we do about it??

— write better code! avoid strcpy and scanf
on untrusted inputs.

(Static analysis is not bad in recent gcc.)

— $W \oplus X$ ("write exclusive-or execute")

Mark call stack as non-executable.

(There's hardware support for this since... a while)
(Observation: attack string would have been on the call stack...)

Adaptation from attackers: return to libc instead...

(Note: libc has tons of potentially useful stuff like `exec(...)`)

addresses of libc functions were actually easy to guess!

Defense to return to libc: ASLR
(Address Space Layout Randomization)

Idea: each process will have library func, stack, and heap in random addresses.

On x86-64, seems to be pretty good entropy!
 $\approx \log_2 1000000$ bits?

— Stack Canaries

Idea: store ^{single} random value in 2 locations:
one on stack between locals + ret. addr:
the other copy is off the stack.

copy of canary ?

ret. addr
...
canary
auth
pwd[0]
:
pwd[0]

Now add to function
epilogue: `if (stack canary \neq copy)
 exit(1);`

Pretty effective!

More sources of vulnerability: printf.

Issue: try to avoid ever using it w/ only one (non-literal) argument!

good: `printf("%s", buffer);`

bad: `printf(buffer);`

dangerous if user supplies buffer!

Say `buffer = "ssssss"`...

This allows one to print data from the stack.

Worse: could actually print memory from an address of the attacker's choice!

Key observation: nasty format string also stored on the call stack and printf is pulling arguments from the call stack:

`"\x7f\xff\xab\x01 %i %i %i %i %i %i %s"`

address you want to read

↑
annoying to set right # of %i's...

`%s` will print what is @ memory location `...`

Scariest still: `%n`. Could write to address `...`!

Further, this could be done without touching any canaries...

Return Oriented Programming (ROP)


Return to libc has some limitations:

— can only call libc functions, might be harder to execute arbitrary code.

Idea: look for useful snippets of assembly (living inside libc code) which end in a 'ret' statement.

Set up memory & registers... then 'ret' to next useful snippet...

Added benefit: can even execute instructions not found in libc!



← accidentally produced i^* as well!

← produced by compiler